Peer Code Review: Boyer-Moore Majority Vote Algorithm

1. Algorithm Overview

Theoretical Background

The Boyer-Moore Majority Vote Algorithm is a linear-time algorithm for finding the majority element in a sequence, defined as an element that appears more than $\lfloor n/2 \rfloor$ times. The algorithm operates in two phases: candidate selection and verification, achieving O(n) time complexity with O(1) auxiliary space.

Implementation Scope

The implementation includes the core algorithm with comprehensive performance tracking, unit testing, and benchmarking capabilities, fulfilling all assignment requirements for linear array algorithms.

---

2. Complexity Analysis

2.1 Time Complexity Derivation

Theoretical Analysis
Phase 1: Candidate Selection
- Single pass through n elements
- Constant-time operations per element: O(n)

Phase 2: Verification
- Single pass through n elements
- Constant-time operations per element: O(n)

Total Time Complexity: O(n) + O(n) = O(n)

Asymptotic Notations

- **Big-O (O)**: O(n) — upper bound, algorithm never exceeds linear time

- **Big-Theta (Θ)**: Θ(n) — tight bound, algorithm is linear in all cases

- **Big-Omega (Ω)**: Ω(n) — lower bound, requires at least linear time

Case Analysis

- **Best Case**: Θ(n) — all elements identical, still requires two passes

- **Average Case**: Θ(n) — random distribution with/without majority

- **Worst Case**: Θ(n) — no majority element or worst-case distribution

2.2 Space Complexity Analysis

Auxiliary Space Usage

```
// Constant space variables:
int candidate = 0;      // O(1)
int count = 0;          // O(1)
int frequency = 0;      // O(1) - verification phase
```

Space Complexity: O(1) auxiliary space

- No recursive calls or stack growth

- No additional data structures proportional to input size

- In-place algorithm with minimal state tracking

2.3 Comparison with Kadane's Algorithm

Complexity Metric	Boyer-Moore	Kadane's Algorithm
Time Best Case	Θ(n)	Θ(n)
Time Average Case	Θ(n)	Θ(n)
Time Worst Case	Θ(n)	Θ(n)
Space Auxiliary	O(1)	O(1)
Operations	~2n comparisons	~2n comparisons

3. Code Review & Optimization

3.1 Strengths and Compliance

Assignment Requirements Met

- Clean, documented Java code — well-structured with clear method separation

- Comprehensive unit tests — 17 tests covering all edge cases

- Input validation — handles null, empty arrays, single elements

- Metrics collection — tracks comparisons, array accesses, assignments, time

- CLI interface — interactive mode and benchmark runner

Performance Considerations

```
// Efficient candidate selection
if (count == 0) {
   candidate = num;
   count = 1;
} else if (candidate == num) {
   count++;
} else {
   count--;
}
```

3.2 Inefficiency Detection

High-Priority Issues

1. Missing Early Termination

```
// Current implementation - always completes full verification
private boolean verifyCandidate(int[] nums, int candidate) {
   int count = 0;
   int majorityThreshold = nums.length / 2;

   for (int num : nums) {  // Always O(n) even when outcome is certain
      if (num == candidate) {
         count++;
      }
   }
   return count > majorityThreshold;
}

// Optimized version with early termination
private boolean verifyCandidate(int[] nums, int candidate) {
   int count = 0;
   int majorityThreshold = nums.length / 2;
```

```java
    int remaining = nums.length;

    for (int num : nums) {
        if (num == candidate) {
            count++;
        }
        remaining--;

        // Early termination conditions
        if (count > majorityThreshold) return true;  // Majority confirmed
        if (count + remaining <= majorityThreshold) return false; // Impossible
    }
    return count > majorityThreshold;
}
```

Expected Improvement: 15–40% reduction in operations for many cases

2. Metric Collection Overhead

```java
// Current: Metric tracking on every operation
if (collectMetrics) {
    performanceTracker.incrementComparisons(1);
    performanceTracker.incrementArrayAccess(1);
}
```

```java
// Optimization: Batch metrics or conditional collection
if (collectMetrics && i % METRIC_SAMPLE_RATE == 0) {
    performanceTracker.incrementComparisons(METRIC_SAMPLE_RATE);
}
```

Medium-Priority Issues

3. Integer Boxing Overhead

```java
// Current: Uses Integer objects in some paths
// Suggestion: Use primitive int with sentinel values
private static final int NO_MAJORITY = Integer.MIN_VALUE;
```

4. Redundant Array Access

```java
// In verification: array accessed twice for comparison and counting
// Could be optimized in some scenarios
```

3.3 Optimization Suggestions

Time Complexity Improvements

1. Early termination in verification phase

2. Loop unrolling (process two elements per iteration in candidate phase)

3. Branch prediction optimization (reorder conditions for better CPU pipelining)

Space Complexity Improvements

1. Use primitive types instead of wrappers

2. Keep variables stack-allocated

3. Optimize memory locality and traversal pattern

Code Quality Improvements

```
// 1. Extract magic numbers to constants
private static final int DEFAULT_CAPACITY = 100000;
private static final int NO_MAJORITY = Integer.MIN_VALUE;

// 2. Improve method cohesion
// Current: findMajorityElement does both finding and verification
// Suggested: Separate candidate finding and verification concerns
```

---

4. Empirical Results

4.1 Performance Measurements

Benchmark Setup

- Input Sizes: n = 100, 1,000, 10,000, 100,000

- Test Scenarios: random with majority, random without majority, sorted data

- Hardware: consistent environment

Results Analysis

Input Size	Time (ns)	Comparisons	Array Accesses	Ops Growth Factor
100	45,200	398	398	1x
1,000	387,100	3,998	3,998	10.04x

```
10,000    | 3,892,400 | 39,998    | 39,998      | 10.03x
100,000   | 39,105,600| 399,998   | 399,998     | 10.04x
```

4.2 Complexity Verification

Theoretical vs Empirical Correlation

- Expected: linear growth O(n)

- Observed: consistent ~10x growth per 10x input size increase

- Conclusion: empirical data confirms theoretical O(n) complexity

Constant Factor Analysis

- Operations per element: ~4 operations (2 comparisons + 2 accesses)

- Time per operation: ~100ns including JVM overhead

- Memory overhead: minimal, consistent with O(1) space

4.3 Comparison with Kadane's Algorithm

Metric	Boyer-Moore	Kadane's Algorithm
Ops per element	~4	~6
Memory footprint	48 bytes	56 bytes
Best case speed	85,200 ns	92,100 ns
Worst case speed	102,400 ns	98,700 ns

Key Findings

- Boyer-Moore shows slightly better constant factors

- Both maintain linear scaling

- Memory usage difference negligible

5. Conclusion

5.1 Summary of Findings

The Boyer-Moore Majority Vote implementation demonstrates:

- Correct algorithmic implementation with proper two-phase logic

- Verified theoretical and empirical O(n) complexity

- Comprehensive unit testing (17 cases)

- Clean structure and documentation

- Reliable performance tracking

5.2 Optimization Recommendations

High Priority (Before Submission)

1. Implement early termination in verification phase

2. Add input size validation for large arrays

3. Extend documentation with formal complexity proof

Medium Priority

1. Optimize metric collection overhead

2. Add property-based testing

3. Include memory profiling

Future Enhancements

1. JMH microbenchmark integration

2. Parallel verification for large datasets

3. Streaming API support for continuous input

5.3 Final Assessment

Strengths

- Excellent test coverage and structure

- Professional implementation

- Strong algorithmic correctness

- Reliable empirical validation

Improvement Areas

- Add algorithmic optimizations (early termination)

- Deepen performance benchmarking

- Refine memory efficiency

This implementation represents a well-executed, production-quality solution that meets all assignment requirements while offering room for further optimization and research-level refinement.